


Data Management Part 1

Yusra AlSayyad
LSST Data Management
Princeton University

DSFP- Session 3
April 25 2017




Why do astronomers care about data management?

- Bigger project teams
- Bigger Datasets
- Data is valuable but time-consuming/costly to manage
- Data Science:
 - **Prepping to run the model:** *"Gathering cleaning matching integrating restructuring transforming loading filtering combining merging verifying extracting" - Bill Howe*
 - Running the model
 - Communicating results
- Goal: recognize what kind of data you have, decide what tools to use

Focus on Data Management Ideas

Google Trends for svn vs git





Today's topics

- SQL Databases in broader context
- Relational Algebra (operations on tables)
- Alternative implementations of Relational Algebra
- Indexes (and Spatial Indexes)

- TOMORROW:
 - Scaling out
 - Parallel databases (including LSST's Qserv), NoSQL databases,
 - Summary of options available and tradeoffs to consider

Databases in Context

Databases fit into a larger context of data management

- How do we store data?



Physically?

Logically?

What would you do
with a million images?

```
--decam
--gapon
--hsc
  --calib
  --raw
  --repo
    --CFHTLS_W1
    --COMET2014F3
    --COSMOS
    --DARK
    --DEEPE01
    --DEEPE02
    --DEEPE03
    --DEEPE04
    --DEEPE05
    --DEEPE06
    --DEEPE07
    --DEEPE08
    --DEEPE09
    --DEEPE10
    --DEN_A
    --DEN_C
    --DEN_E
    --DITH 14H
```

What would you do
with a million images?

```
--decam
--gapon
--hsc
  --calib
  --raw
  --repo
    --CFHTLS_W1
    --COMET2014F3
    --COSMOS
      --2014-02-02
      --2014-02-05
      --2014-03-26
      --2014-03-27
      --2015-01-17
      --2015-01-18
      --2015-01-19
      --2015-01-20
    --DARK
    --DEEPE01
    --DEEPE02
    --DEEPE03
    --DEEPE04
    --DEEPE05
    --DEEPE06
```


Time to count all images
taken:

With HSC

in the COSMOS field,
on March 26th 2014:

```
$ time ls  
hsc/repo/COSMOS/2014-  
03-26/00815/HSC-Y |  
wc -l
```

2912

```
real    0m0.039s  
user    0m0.029s  
sys     0m0.008s
```

```
--decam  
--gapon  
--hsc  
  --calib  
  --raw  
  --repo  
    --CFHTLS_W1  
    --COMET2014F3  
    --COSMOS  
      --2014-02-02  
      --2014-02-05  
      --2014-03-26  
        --00815  
          --HSC-Y  
            --HSC-0000628-000.fit  
            --HSC-0000636-038.fit  
            --HSC-0000662-042.fit  
            --HSC-0000670-080.fit  
            -- (and 2912 more files)  
      --2014-03-27  
      --2015-01-17  
      --2015-01-18  
      --2015-01-19  
      --2015-01-20
```

Data Models are Unavoidable

All images taken with the HSC-Y band?

```
$ time ls
hsc/repo/*/*/*/HSC-
Y | wc -l
193171
```

```
Real    48m8.737s
user    0m5.824s
sys     1m2.072s
```

```
--decam
--gapon
--hsc
  |--calib
  |--raw
  |--repo
    |--CFHTLS_W1
    |--COMET2014F3
    |--COSMOS
      |--2014-02-02
      |--2014-02-05
      |--2014-03-26
      |--00815
        |--HSC-Y
          |--HSC-0000628-000.fit
          |--HSC-0000636-038.fit
          |--HSC-0000662-042.fit
          |--HSC-0000670-080.fit
          |-- (and 2912 more files)
      |--2014-03-27
      |--2015-01-17
      |--2015-01-18
      |--2015-01-19
      |--2015-01-20
```

Table (or Relational) Structure

Id	taiObs	expId	point	dataType	visit	dateObs	fra
1	2014-02-02	HSCA90585000	763	OBJECT	905850	2014-02-02	HSO
2	2014-02-05	HSCA90794800	766	OBJECT	907948	2014-02-05	HSO
4	2014-02-03	HSCA90679200	764	OBJECT	906792	2014-02-03	HSO
5	2014-02-01	HSCA90554800	762	OBJECT	905548	2014-02-01	HSO
7	2014-02-04	HSCA90751400	765	OBJECT	907514	2014-02-04	HSO
9	2014-02-05	HSCA90806200	766	OBJECT	908062	2014-02-05	HSO
10	2014-02-03	HSCA90699800	764	OBJECT	906998	2014-02-03	HSO
11	2014-02-01	HSCA90560000	762	OBJECT	905600	2014-02-01	HSO
12	2014-02-05	HSCA90794000	766	OBJECT	907940	2014-02-05	HSO
13	2014-02-01	HSCA90565800	762	OBJECT	905658	2014-02-01	HSO
14	2014-02-01	HSCA90542000	762	OBJECT	905420	2014-02-01	HSO
15	2014-02-01	HSCA90563200	762	OBJECT	905632	2014-02-01	HSO
16	2014-02-03	HSCA90655200	764	OBJECT	906552	2014-02-03	HSO
17	2014-02-04	HSCA90761200	765	OBJECT	907612	2014-02-04	HSO
19	2014-02-03	HSCA90664400	764	OBJECT	906644	2014-02-03	HSO
20	2014-02-03	HSCA90702600	764	OBJECT	907026	2014-02-03	HSO
Etc...							

Table (or Relational) Model Example

Source Table

sourceld int NOT NULL,
ra float,
dec float,
mag float,
exposureId, bigint,
PRIMARY KEY (sourceld)

Exposure Table

exposureId bigint NOT NULL,
fwhm float,
fluxMag0 float,
filter float,
Mjd int,
path varchar(max),
PRIMARY KEY (exposureId)



Table (or Relational) Structure

SourceId	psfMag	filter	fluxMag0
1	21.1	2	5721049381928.0
2	20.1	2	5721049381928.0
4	19.3	2	5721049381928.0
5	18.3	2	5721049381928.0
7	19.6	2	5721049381928.0
9	20.1	2	5721049381928.0
10	20.1	2	5721049381928.0

```
SELECT s.sourceld, s.psfMag, e.filter, e.fluxMag0
FROM Source s
INNER JOIN Exposure e
ON s.exposureId = e.exposureId
```

How many friends does Alice have?

Friendships

PersonId1
PersonId2
BeginDate

Person

PersonId
FirstName
LastName

```
SELECT a.Person2  
FROM Friendships a  
INNER JOIN Person p  
ON p.personId = a.personId1  
WHERE p.FirstName = 'Alice'
```

UNION ALL

```
SELECT a.Person1  
FROM Friendships b  
INNER JOIN Person p  
ON p.personId = a.personId2  
WHERE p.FirstName = 'Alice'
```

How many *friends of friends* does Alice have??

Friendships

PersonId1
PersonId2
BeginDate

Person

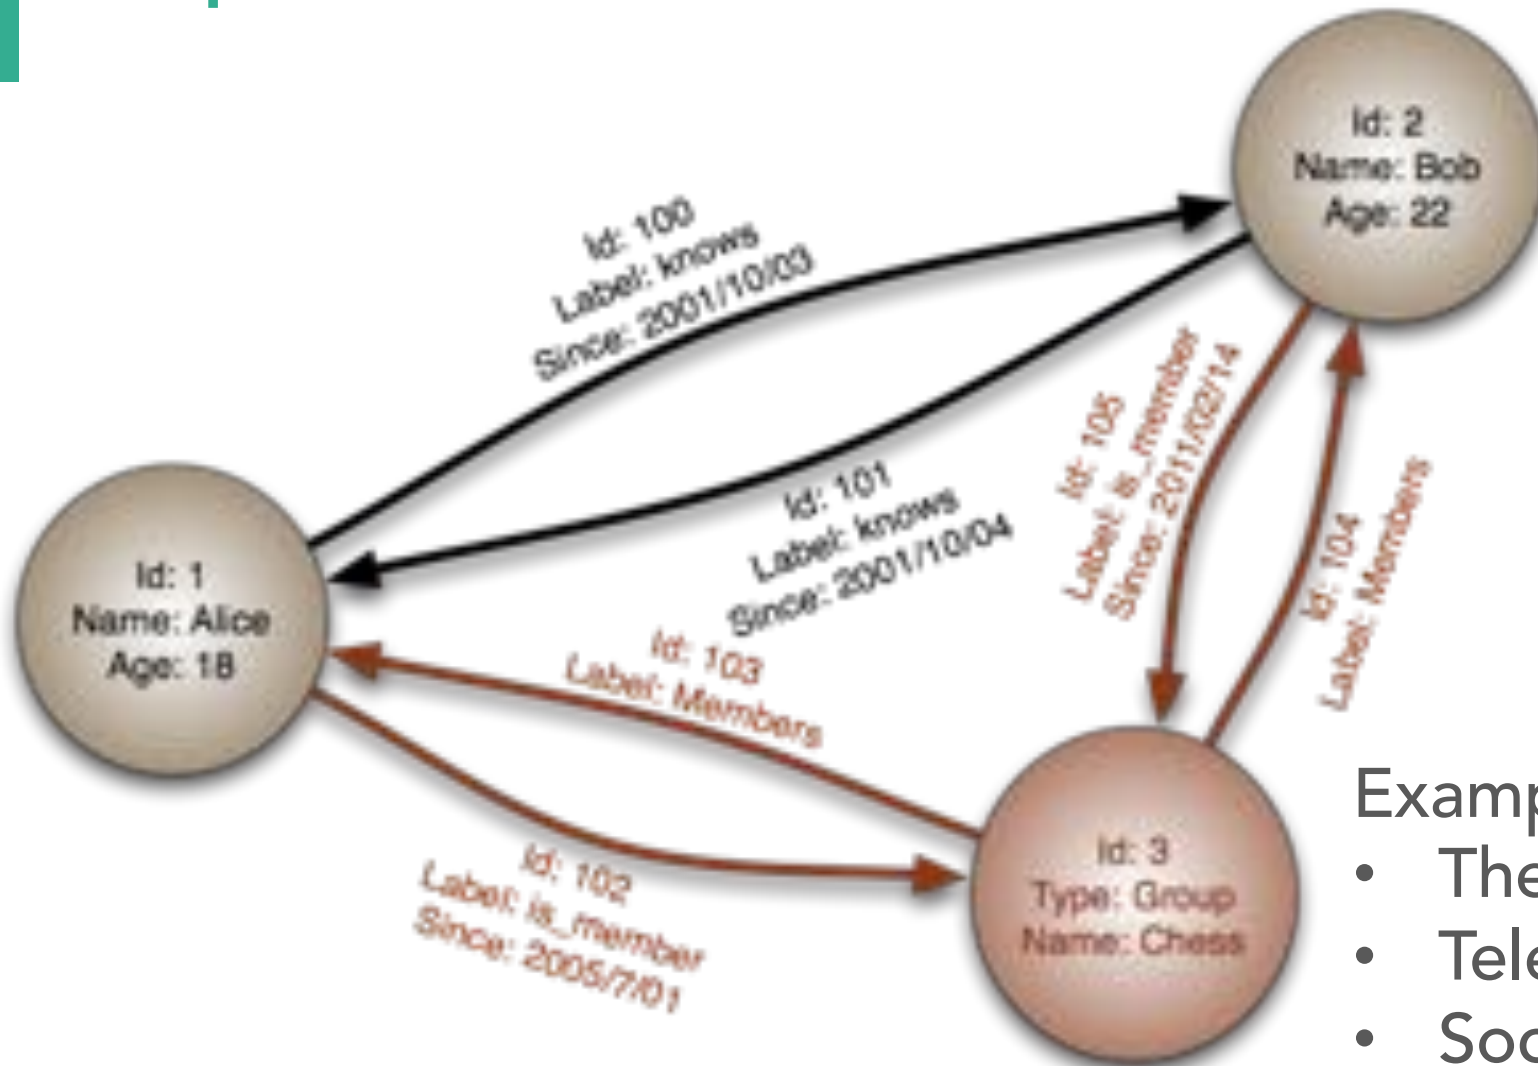
PersonId
FirstName
LastName

```
SELECT a.Person2  
FROM Friendships a  
INNER JOIN Person p  
ON p.personId = a.personId1  
WHERE p.FirstName = 'Alice'
```

UNION ALL

```
SELECT a.Person1  
FROM Friendships b  
INNER JOIN Person p  
ON p.personId = a.personId2  
WHERE p.FirstName = 'Alice'
```

Graph Structure



Examples:

- The web
- Telecom
- Social Networks

Image credit: wikipedia

Data models are defined by

■ Structures

- Tree
- Relational (Tables)
- Graphs (*nodes and edges*)
- Keys and values

■ Constraints

- Examples:
 - File can't be in two folders.
 - All rows have same number of columnnd

■ Operations

SQL as Relational Algebra

Operations on table implemented by Relational Database Management Systems (RDMS)

Relational Algebra

- Basic operations:
 - *Selection* (σ) Selects a subset of rows from relation.
 - *Projection* (π) Deletes unwanted columns from relation.
 - *Cross-product* (\times) Allows us to combine two relations.
 - Inner join, left join, right join, outer join
 - *Set-difference* ($-$) Tuples in A, but not in B
 - *Union* (\cup) Tuples in A or B
- Additional operations:
 - Intersection, *join*, division, renaming
- Since each operation returns a relation, operations can be *composed!* (Algebra is “closed”.)

Adapted from Database Management Systems, R. Ramakrishnan and J. Gehrke



Logical Join Types

Aggregate Operators and GROUP BY


- COUNT (*)
- COUNT ([DISTINCT] A)
- SUM (A)
- AVG (A)
- MAX (A)
- MIN (A)

```
SELECT plate, COUNT(id)
FROM SpecObjAll
GROUP BY plate
```

Alternative implementations of Relational Algebra




Interlude to talk about some tools

SQL Databases are not necessary to get Relational Algebra



Wes McKinney's personal story on birth of pandas:
youtube.com/watch?v=kHdkFyGCxiY

pandas
 $y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$



You can **SOMETIMES** use pandas if the dataset doesn't fit into memory

```
dc = pd.read_csv(filenameDeepCoadd, compression='gzip',
                  header=None, index_col=0)



sources = pd.read_csv(filenameDeepSource,
                      compression='gzip', header=None,
                      chunksize=100000, index_col=0)

for i, chunk in enumerate(sources):
    # Now you can join the chunk after some prep:
    joined = pd.merge(chunk, dc, left_on='deepCoaddId',
                      right_index=True)

    # And so on
```

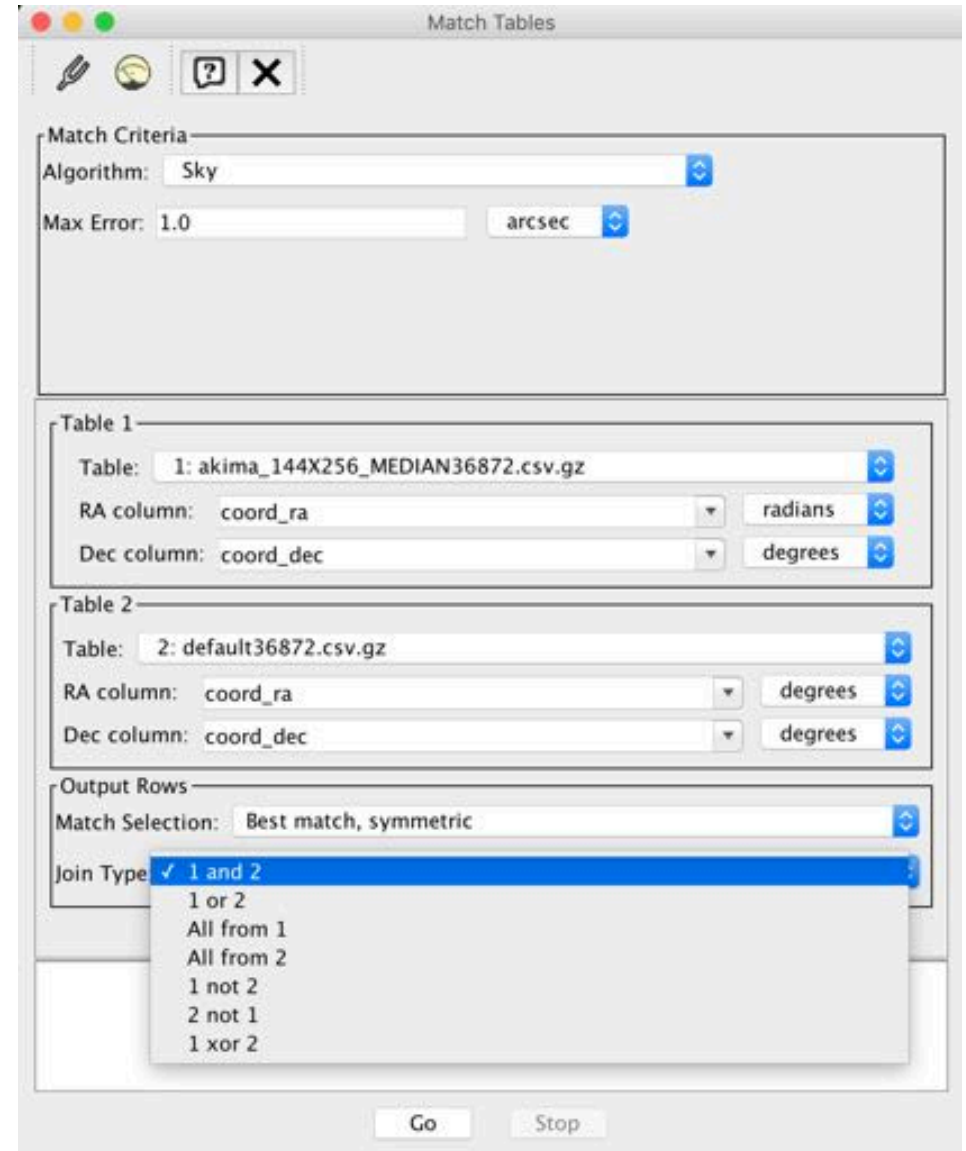

Beware of over-engineering

"I SPEND A LOT OF TIME ON THIS TASK.
I SHOULD WRITE A PROGRAM AUTOMATING IT!"



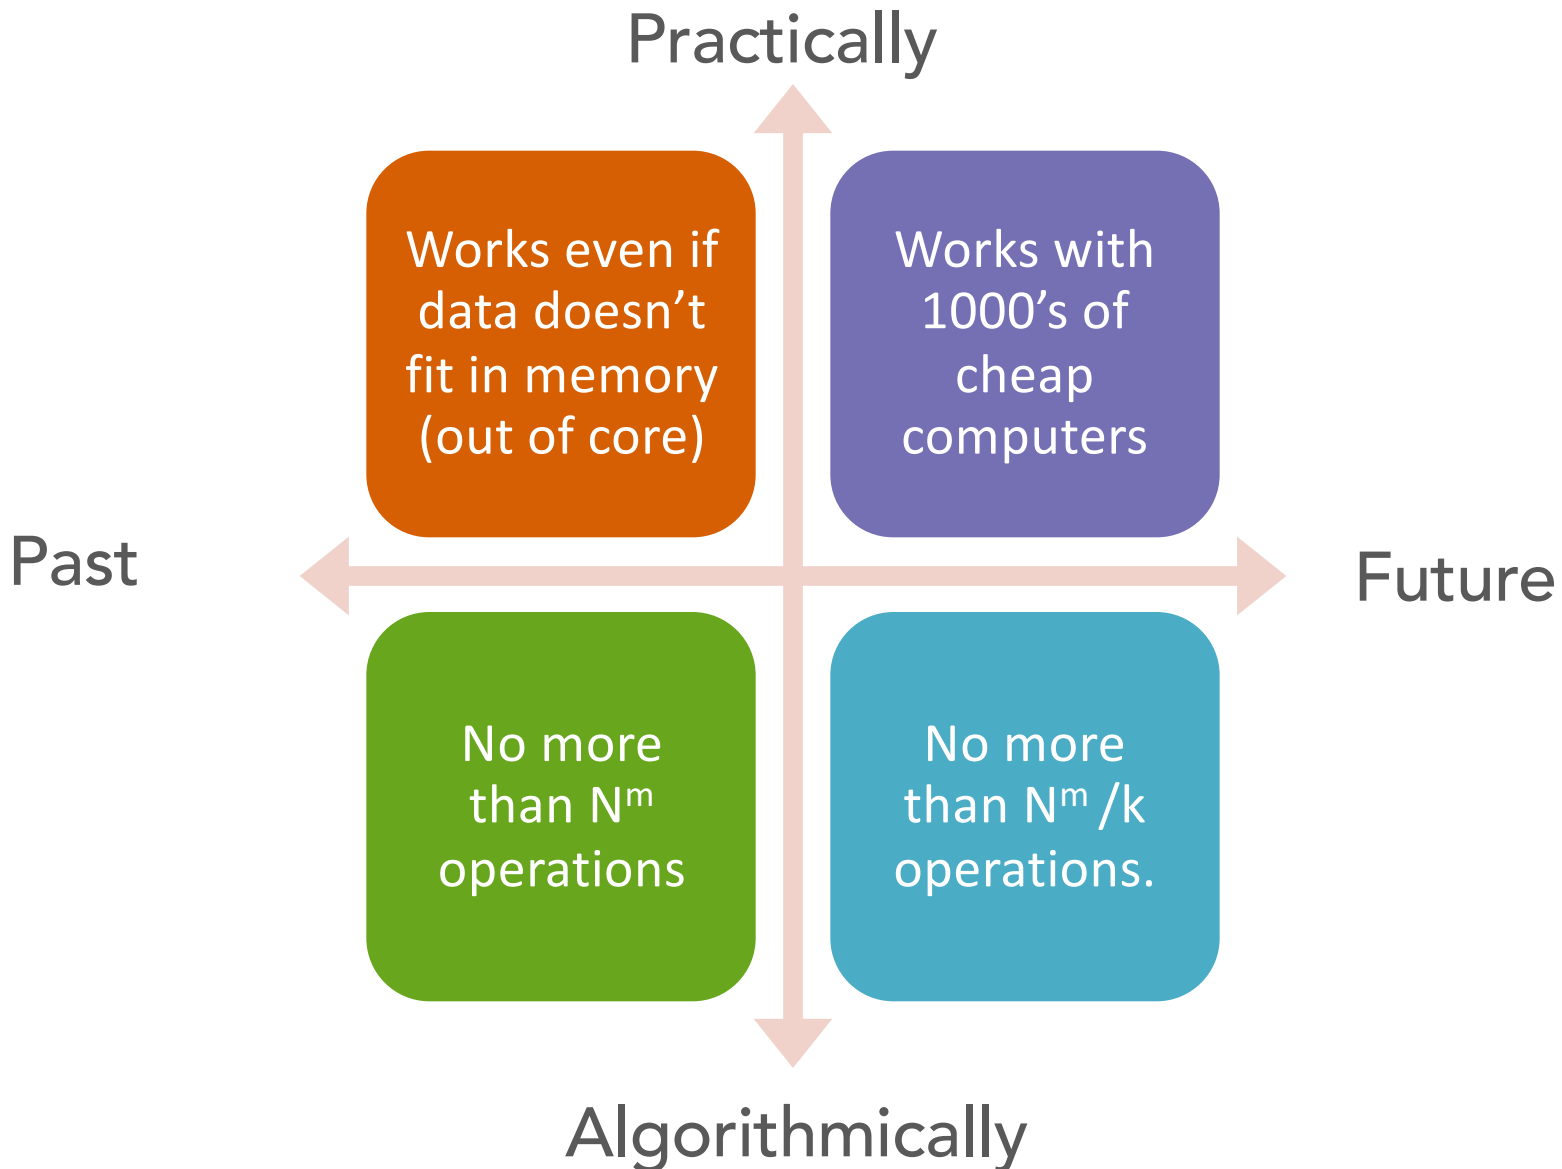
TOPCAT is written specifically for astronomers by an astronomer

GUI tool for relational algebra on tables
+ fits table I/O
+ theta joins ON RA/DEC



Indexes (and Spatial Indexes)

What do we mean by scalability?



Scalability

- Practically:
 - Before: Works even if data doesn't fit in memory (out of core)
 - Now: Can we utilize 1000's of cheap computers
- Algorithmically :
 - For N data items:
 - Before: You can do no more than N^m operations
 - Now: You can do no more than N^m / k operations
 - Soon: You can do no more than $N \cdot \log(N)$ operations (one pass – streaming)

Example: search for records where value = 41

$N = 20$

How many comparisons must we make?

41	72	0	30	14	9	18	34	39	53	41	68	20	87	2	67	41	55	14	19
----	----	---	----	----	---	----	----	----	----	----	----	----	----	---	----	----	----	----	----

Example: search for records where value = 41

$N = 20$

How many comparisons must we make?

0	2	9	14	14	18	19	20	30	34	39	41	41	41	53	55	67	68	72	87
---	---	---	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----	----

B+ Tree Index

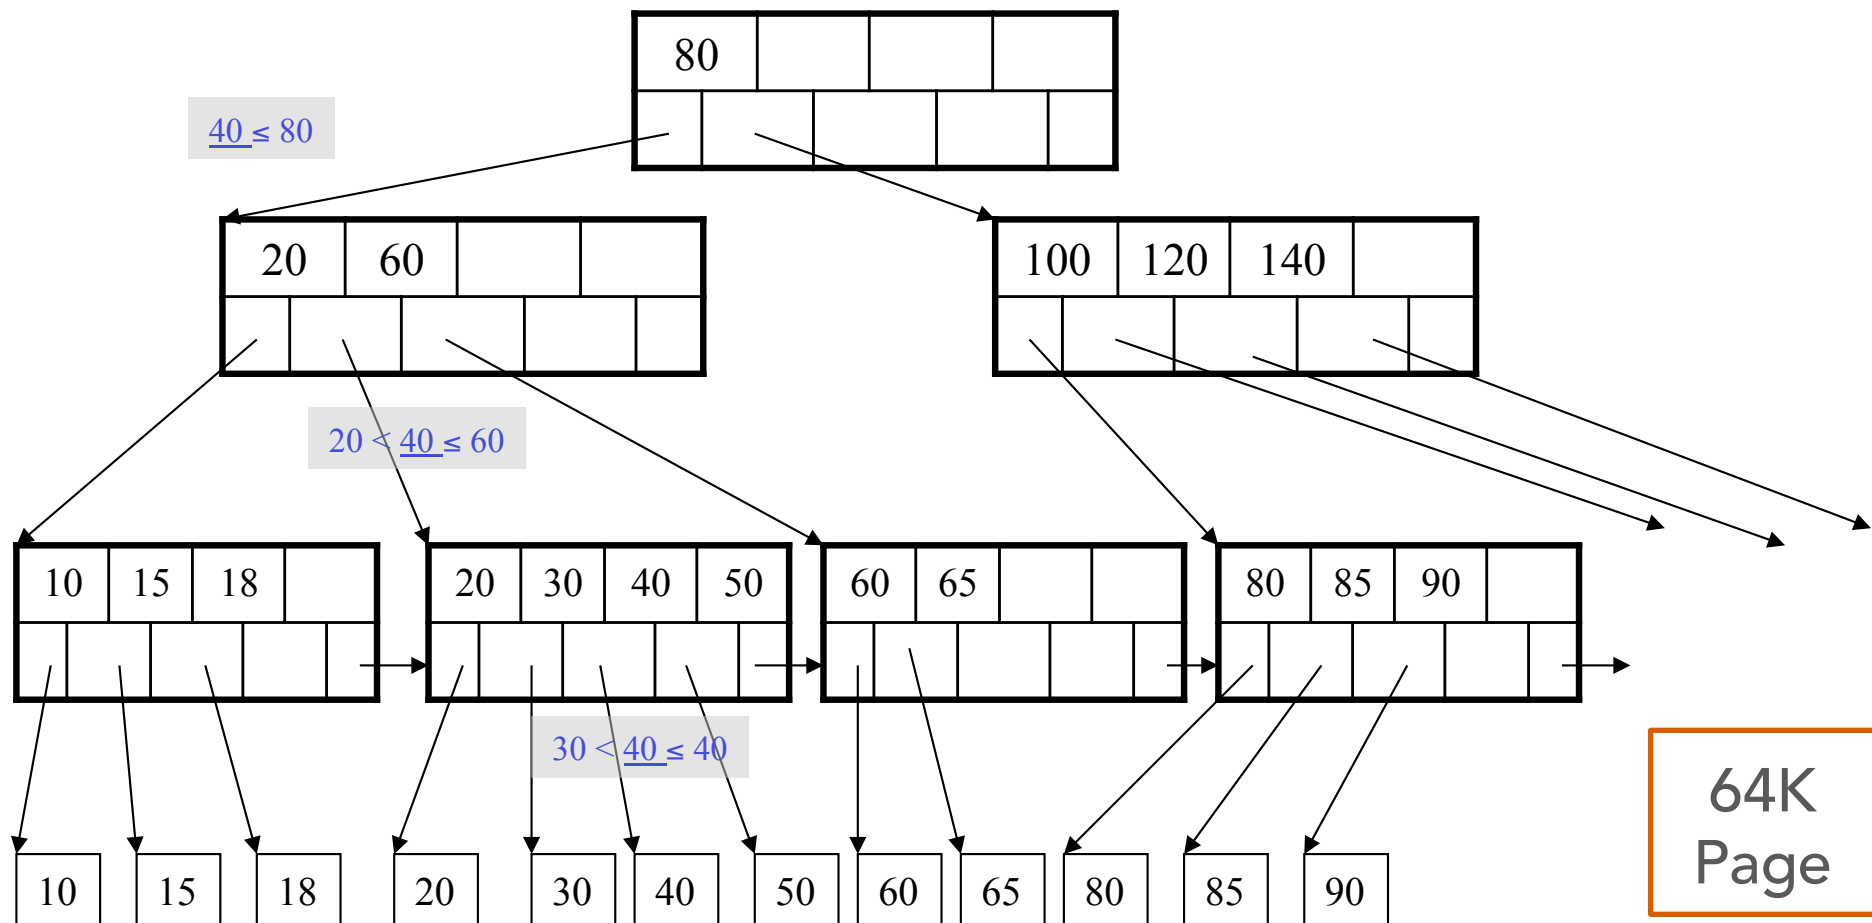


Image credit: UW CSE

Relational databases excel at finding record sets within large datasets



What are good choices of columns for indexing?

Source Table

sourceld int NOT NULL,
ra float,
dec float,
mag float,
exposureId, bigint,
PRIMARY KEY (sourceld)

Exposure Table

exposureId bigint NOT NULL,
fwhm float,
fluxMag0 float,
filter float,
Mjd int,
path varchar(max),
PRIMARY KEY (exposureId)



```
CREATE INDEX idx_exposureId ON Source(exposureId);
```

In astronomical applications those seeks are often spatial





But coordinates are usually stored as RA, DEC

What do we index on

Many databases ship with Geographic Information System (GIS) support or have plugins

- Select a spherical model of the Earth
- Map:
 - R.A. --> Longitude
 - Dec. --> Latitude




OpenGIS®

See en.wikipedia.org/wiki/Spatial_database for a long list

Astronomers are fond of Hierarchical Triangular Mesh



Reduce 2 dimensions to 1 with space filling curve



Background Information for the challenge problem

